

Resume Change Request — Draft 1 Feedback

To: Jaya Waldia, Resume Writing Team, Naukri

From: Sameer Khan

Ref: #VR#260817TS43026084_43453892#

Re: CustCopy2.pdf / CustomerCopy.docx , received 21 Aug 2026

Attached: resume_sameer_khan_v2.pdf — a reference layout I have prepared showing all of the changes below already applied. It is a reference for the wording and ordering, not a replacement for your template; please treat it as the target content.

Thank you for the draft — the structure is a clear improvement on what I sent you, and several of your decisions are better than mine (listed in §8). I have gone through it line by line and prepared this document so my feedback is specific and testable rather than a matter of taste.

Requests are ranked. **§1 is a factual correction and is not negotiable.** §2–§6 are strongly evidenced. §7 is a question, not a request.

1. Factual error — the F1 figure reads backwards (*must fix*)

Your draft, ICE bullet 4:

*"Improved model performance across 946K+ pages by identifying a train-to-benchmark taxonomy mismatch **that had reduced F1 from 0.62 to 0.92**"*

"Reduced from 0.62 to 0.92" describes an *increase*. As written, the sentence says the defect caused the improvement.

What actually happened: the taxonomy mismatch had depressed F1 **to** 0.62. After I fixed it, F1 rose **to** 0.92.

Please use:

*"Recovered a production accuracy regression across 946,355 pages by fixing a train/benchmark taxonomy mismatch **that took F1 from 0.62 to 0.92**; further established that 64% of residual errors were valid predictions discarded by a confidence threshold, leading to calibrated per-class selective prediction across all 643 classes."*

Why this matters more than the wording: every interviewer for these roles is an ML practitioner. They will notice within one read. The risk is not the typo — it is that it implies I did not check my own numbers.

2. My current role is written in the past tense

ICE Data Services is dated **Nov 2023 – Present**, but every bullet reads Owned , Managed , Modernized , Developed . On a current role this reads as though I have left.

Please convert the ICE bullets to present tense — Own , Manage , Lead — and keep past tense only for completed one-off events (the P1 RCA, the TF1→TF2 migration).

Prior roles stay in past tense, correctly.

3. Core Competencies — please remove, and use the space for domains

I want to be precise here, because your cover note gives a specific rationale for this block: *"keywords mentioned in Core Competencies and across Organisational Experience have greater chances of finding a match in Applicant Tracking Systems."*

I tested that claim against your own document. I extracted every content word from the 12 competency phrases and searched for each one in the rest of the resume.

Competency phrase	Its keywords, counted elsewhere in the same resume
Document AI & Intelligent Document Processing	document×24, processing×9 — <i>only new word: "intelligent"</i>
Machine Learning & Multimodal Deep Learning	machine×4, learning×8, multimodal×5, deep×4
Computer Vision & OCR-Free AI	computer×4, vision×5, ocr-free×3
NLP & Information Extraction	nlp×1, information×2, extraction×4
LLM Engineering & Generative AI	llm×5, engineering×4, generative×4
RAG & Enterprise Retrieval Systems	rag×6, enterprise×2, retrieval×7, systems×3
Transformer Architecture & Model Fine-Tuning	transformer×5, architecture×7, model×22, fine-tuning×1
Distributed ML & Large-Scale Model Training	distributed×5, large-scale×2, model×22, training×9
MLOps & Production AI Engineering	mlops×4, production×18, engineering×4
AWS SageMaker & Cloud AI Engineering	aws×6, sagemaker×8, cloud×3, engineering×4
Model Optimization, Inference & AI Reliability	model×22, optimization×4, inference×2, reliability×2
AI Solution Architecture & Technical Leadership	solution×5, architecture×7, technical×3 — <i>only new word: "leadership"</i>

Result: 10 of the 12 phrases contribute no keyword that is not already in the document, usually several times over, and mostly in the Technical Skills block immediately below them. The only two new words the block adds are **"intelligent"** and **"leadership"** — neither of which a recruiter searches for.

There is also an extraction problem

I pulled the text layer out of `CustCopy2.pdf` — this is what a parser receives:

```
CORE COMPETENCIES
EDUCATION
  B.Tech. – Computer Science & Engineering...
CERTIFICATIONS
PERSONAL DETAILS
CAREER TIMELINE
Document AI & Intelligent Document Processing
Machine Learning & Multimodal Deep Learning
...
```

The `CORE COMPETENCIES` heading is separated from its own items, which land after `CAREER TIMELINE` . A keyword-matching ATS still finds the strings, but any parser that maps content to sections files them under the wrong heading.

The same extraction also yields `Model FineTuning` — the hyphen is destroyed by the line break inside the box. A search for "fine-tuning" or "fine tuning" will not match it.

One competency I specifically want removed regardless

"AI Solution Architecture & Technical Leadership." I have never held a lead or manager title. If a resume claims technical leadership, the interview will test it, and I would rather compete on what I can evidence.

What to put in that space instead

Domain Expertise — three items, none of which appear anywhere else in the document:

- Mortgage & Loan File Documents
- Retail Banking — Arabic / English
- Identity Verification & KYC

For Document AI roles, domain match is frequently what shortlists a candidate, and this is genuinely absent from the current draft.

If the block must stay for template reasons: please trim it to 5–6 items, drop the Technical Leadership entry, and build it as a single-column table with cell shading rather than text boxes or SmartArt — drawing objects sit outside the document flow and many parsers skip them entirely, whereas table cells are extracted reliably.

4. Please restore seven specific phrases from my original

These are not stylistic preferences. Each one is a question I want an interviewer to ask me, because each is something I can talk about for ten minutes. Generalising them removes the hook.

My original	Your draft	What is lost
"a multi-head-attention Transformer I wrote from scratch"	"custom multi-head-attention Transformer components"	Personal authorship. This single phrase generates more technical interviews than anything else on the page.
"Proposed and built the OCR-free successor"	"Developed an OCR-free Donut/Swin model"	Initiative. "Proposed" says I identified the opportunity; "developed" says I was assigned it. For a senior IC role that distinction is the whole point.
"Led the P1 root-cause analysis that cleared the model"	"Led technical RCA for a P1 production regression"	The outcome. Without it the bullet describes activity, not result — and the result was that my model was exonerated.
"an input validator was rejecting whole documents, so one dropped page could null a 530-page file"	"rejected complete documents when individual pages failed"	The concrete image. "530-page file" is what a reader remembers an hour later.
"I ran the searches inside label and page-count groups so I could use many small exact searches instead of one approximate index"	"Optimized exact retrieval through label/page-count grouping"	The engineering judgement. At senior level the reasoning is the signal; the mechanism alone reads as a tool list.
"the bank issued it as a circular... That fixed the problem at the source instead of in preprocessing"	"established a standardized configuration subsequently mandated bank-wide"	The systems-thinking punchline — I fixed a data problem upstream instead of patching it in code.
"9 years in Document AI... applied to one problem across three roles"	(absent)	My single most differentiating sentence. Nine years of depth on one problem is rare; the draft replaces it with breadth claims that every candidate makes.

5. Hard numbers appear too late

Counting words from the top of page 1 to the first concrete metric:

Version	Words before the first hard number
Your draft	~370 (Job Objective + 6 Profile Summary bullets + competency boxes)
My original	~30

A recruiter's first pass is short. In the current draft it lands on "architecture-oriented," "end-to-end exposure," "robust model architectures" — phrasing that does not distinguish me from any other applicant. **"643 classes at 1.7M requests/day with 99.99% success"** does.

Request: compress the Job Objective to two lines and the Profile Summary from six bullets to three, so a hard number appears in the first two lines of the page.

I am not asking to shorten the resume overall — your Organisational Experience section is a good length. The density is concentrated entirely in the top block.

6. Three hedging phrases to remove

Current text	Problem
"with knowledge and practical exposure to LLM-based applications"	"Exposure" is read by senior hiring managers as <i>has not actually done this</i> .
"with growing focus on Agentic AI and modern GenAI architectures"	Signals aspiration, not capability — and there is no agentic work on the resume to support it.
"with a strong understanding of designing scalable ML pipelines"	"Understanding" is weaker than the bullets below it, which describe systems I actually built and run.

Request: state these concretely or cut them. I would rather claim less and defend all of it.

7. Question, not a request — has this template been parser-tested?

The two-column layout with a left sidebar is visually strong. My only concern is that some enterprise ATS platforms interleave sidebar text into the body when extracting.

Has this specific template been tested against Workday, Greenhouse and Taleo parsers, or against Naukri's own parser? I am not assuming there is a problem — I would just like to know, because most of my applications will go through company portals rather than the Naukri portal.

If it has not been tested, I am happy to keep two versions: your designed layout for the Naukri portal and recruiter email, and a single-column variant for strict ATS portals. That is a good outcome, not a complaint.

8. The blue fields you asked me to complete

Field	My answer
Certifications	I have none. Please remove the section entirely rather than leaving it blank — an empty heading draws attention to the gap.
Date of Birth	Please remove this field. Several of my target employers are US/UK-headquartered, where DOB on a resume is discouraged and occasionally triggers compliance filtering. It adds nothing for domestic applications either.
Languages Known	English, Hindi, Telugu, Urdu (<i>please confirm this list is correct before it goes in</i>)
"over 9 years of experience"	Correct. Feb 2017 – present.

9. What I want kept from your draft

To be clear that this is not a rejection — these are improvements on what I sent you, and I want them preserved:

1. **Technical Skills keyword coverage.** Your categorisation is better than mine, and you added terms I had missed — Information Extraction, Model Optimization, Production Monitoring, Confidence Calibration, Image Preprocessing, Instruction Tuning. Keep all of it.
2. **Deduplication and signature verification promoted to page 1.** In my version both were buried on page 2. You were right to surface them.
3. **The Zyclyx role split into separate bullets.** I had crammed classifier, detection, OCR, NER and signature verification into one dense bullet. Yours is far more readable.
4. **The individual-contributor targeting statement.** Useful — it pre-empts the "do you want to move into management?" conversation. I only want it shorter, not gone.
5. **Reverse-chronological structure and the overall two-page discipline.**

10. Questions for you

1. Has this template been tested against non-Naukri ATS parsers (§7)?
2. Can you send the **editable .docx** with each revision, so I can check wording without a round trip?
3. If the Core Competencies block is a fixed requirement of the service, can it be built as a shaded table rather than text boxes (§3)?
4. Is there a length constraint I should work within if I want the top block compressed (§5)?

Appendix — how the evidence in §3 and §5 was produced

- **Keyword counts:** every content word was extracted from the 12 competency phrases, stop-words removed, then counted with word-boundary matching against the full text of `CustCopy2.pdf` excluding the competency block itself.
- **Extraction order and the FineTuning artifact:** taken from the PDF's own text layer — the same layer an ATS reads. Reproducible from the file you sent.
- **Word-position counts in §5:** counted from the first word of the document to the first numeric metric.

Happy to walk through any of this on a call.

